COS30031 Games Programming
Custom Project Report

Daniel Coady 102084174

04/11/2021

Contents

1 Introduction

1.1 Background

Computers are incredible and complicated these days, with many components that allow for the more efficient and effective computation of various tasks. For me personally, by far the most interesting modern day development (which many will actually take for granted) is the humble GPU. Short for graphics processing unit, it is a core aspect to any computer whether it is pushing hyper-realistic game graphics or displaying the words on your social media feeds. These GPUs differ greatly from the CPU which many view as the beating heart of any computer. CPUs have been designed to be more general purpose, allowing for operations such as logical branching, maths, and bitwise magic. GPUs on the other hand are entirely focused on floating point maths operations, and have been architected to be able to do those rapidly and in parallel. This means modern day graphics requirements, such as being able to independently address one of the literal millions of pixels on your screen (a standard 1080p screen contains a little over 2 million pixels!), is now a trivial task that can be completed with great speed.

Computers haven't always been like this however, and many moons ago we would have required the CPU to perform all of the logical and graphical work of a computer. This complicated matters greatly as soon as you wanted to display complex graphics, and it only got worse if you wanted to display complex graphics *and* provide complex logic–such as in a game. It's for this reason that many games of yore have had to come up with some very clever techniques to "cheat" graphics. Most games for the longest time were purely 2D since that was about as much as we could reasonably handle with the hardware at the time. Before 3D accelerated technology hit the general consumer market, 3D games were but a dream... Kind of.

Enter the raycaster. There will be more later on about how it works and why it's so fast, but for now just know that raycasters were some of the earliest attempts at creating 3D graphics in games. Not to be confused with ray marching and ray tracing (two very cool 3D graphics techniques as well!), raycasters have a very distinct look to them that many will remember from the first Wolfenstein 3D game. They tend to look very blocky, with billboarded sprites representing the entities within a scene. This technology forms the basis of my custom project.

1.2 The Project

At a high level, this project will take form of a first person shooter using 3D raycaster graphics. The front end will use SFML to display the internal frame buffer as well as receive/process the window events. The back end will be far more complex, consisting of the actual raycasting engine, an ECS implementation for managing enemy entities, multi-threading to assist with performance, and some simple collision detection. The entire project has been written in C++

within a Linux environment using the gcc toolchain, and borrows greatly both from Austin Morlan's writeup on ECS[1] and Dmitry V. Sokolov's tinyraycaster series of tutorials[2].

2   Implementation

2.1   Front End

The bulk of the front end in this project is handled by the Simple and Fast Multimedia Library, or SFML for short. SFML is similar to SDL in a variety of ways–both have many useful abstractions of low level graphics, image handling, audio processing, etc. There are some key differences between the two however:

- SDL has a very C style design and API, while SFML's design and API are much more like C++

- SFML's abstractions tend to be easier to work with due to it's object oriented nature

- SDL's graphics abstractions are fairly limited, only allowing for simple quads and textures while SFML has more primitives and a simple shader pipeline

- SFML just generally has less boilerplate code than SDL

All of the above ultimately influenced my decision to choose SFML over SDL as we had been using for all of our previous tasks throughout the unit. Add to this that it has a very sensible system for event polling that is similar (though in my opinion, slightly better) to SDL's own event system, and it just made sense for the purposes of my project.

There is another option however, one that I think might have made more sense for this project actually: raw OpenGL. It's a scary prospect, for sure, but there's a big benefit in the form of performance that we would have gained. There is of course some added complexity that would come from this since we're essentially working with little to no abstractions. However, considering that there isn't much we need from SFML outside of the events system and drawing textures to a quad, it would have been decently trivial to make the change over to raw OpenGL.

2.2   Back End

2.2.1   Enemy AI

Enemies are all treated as entities in an Entity Component System, or ECS for short. This is quite a simplified version of ECS, as it only need serve one specific

---

[1]https://austinmorlan.com/posts/entity_component_system/
[2]https://github.com/ssloy/tinyraycaster/wiki/Part-0:-getting-started

purpose with little to no generics involved. This has allowed me to keep it as simple as it is, while also ensuring maximum performance from both algorithmic and cache optimisations.

Each entity will contain a component for movement, the enemy type, and the distance it is from the player. They have been separated like this since there are situations where not all of the data related to an entity would be necessary, and so by lowering the amount of data we need to retrieve in each of the related systems we can help ensure higher amounts of data are being fetched. In terms of systems present, there are a few present. The main ones are the render system, and the enemy behaviour system. The render system is fairly uninteresting as all it does is read the components for each entity to render them accordingly. The enemy behaviour system on the other hand is far, far more interesting. This system, as the name would imply, is where all of the actual AI for each enemy lives.

The enemy AI is incredibly simple, as it exhibits only two behaviours:

- Chasing the player

- Waiting to see the player

Chasing the player is fairly simple: calculate the unit vector pointing towards the player, then add that vector to the current position scaling by speed and current delta time. This only happens when the enemy can see the player however, and this was a fair bit more complex. You see, physics raycasting with my current solution is rather complicated due to there being no real collision detection to speak of. If I had implemented some simple collision detection with axis-aligned bounding boxes then I may have been able to do some simple maths to check for ray intersections with edges of the AABB, but there is no such thing implemented. So there's only really one option left for me: a sort of ray march. To be clear, this isn't the graphics technique for volumetric raycasting, I just don't have a better name for it at the time of writing this. Simply put, the way that the physics raycasting works is by continually "marching" in a given direction until it hits something. If we hit something, then we register what that is and make a decision based on that. In our case, if we hit a wall then we don't want to move the entity. If it hits the player, then we start chasing.

This gets expensive to perform every frame however, so I wanted a way to make it more performant. This posed an interesting issue since more performance seemed to require me to give up some amount of accuracy. This lead me to perform a series of tests on how performant and accurate different levels of coarseness were when performing a physics raycast, but perhaps most importantly I wanted to find out at what point did the believability of the AI suffer. Much to my surprise, I could go incredibly coarse before it started to show any sort of behaviour that was odd. And so, I made the raycast treat each map tile as a single unit, and it will draw a line from enemy to player checking for a hit. This hit a good balance, where the entities never ended up showing strange behaviour, but performance is hardly impacted due to how coarse the physics raycast calculations are.

2.2.2   Collision Detection

Unfortunately, there was not much in the way of physics collision detection implemented into this project. Much of this comes down to problems with time constraints, as there has not been enough time to work it in, and the limited time I have had at my disposal caused me to make some poor decisions when it came to architecture which ended up not lending itself well to implementing even simple collision systems.

Bar the previously mentioned physics raycasting for enemy movement, there is one other area where I have implemented collision detection: player movement. Well... In theory at least. The idea is that the player shouldn't be able to walk into walls (both because walls are solid, but also because it crashes the game) and this is achieved with the following code:

```
// in the update method
if ( get_map_tile( int(player.position.x), int(old_pos.y) ) != Floor ) {
    float wall_start = floor( player.position.x );
    if ( move_vec.x < float(0) ) wall_start += 1.0;
    else wall_start -= 0.05;
    player.position.x += wall_start - player.position.x;
}

if ( get_map_tile( int(old_pos.x), int(player.position.y) ) != Floor ) {
    float wall_start = floor( player.position.y );
    if ( move_vec.y < 0 ) wall_start += 1.0;
    else wall_start -= 0.05;
    player.position.y += wall_start - player.position.y;
}
```

The idea is that when we move, we want to check if there's a wall where we have just tried to move. If there is, then we move the player back a small amount to ensure they don't go inside of the wall. Now this actually works excellently, but only without multi-threading. This is something that I have tried to resolve and have failed, even with my attempts to ensure that all resources in the hot path are locked with a mutex. It's really quite strange, but once again due to time constraints I was unable to do much more with it and had to cut my losses.

2.2.3   Multi-threading

Multi-threading is something that you see thrown around a lot as a way to get free or easy performance. Well, I'm happy to report that those people are big fat liars. There are many more complications that come along with implementing multi-threading into any application, and I encountered a fair few of these while working it into my current solution.

For some context, the goal of this was to completely decouple the front end from the back end in such a way that they don't rely on each other executing in sequence. This allows for one or the other to complete their own loops without

being blocked by the other, which is incredibly important for the front end to ensure that the application always feels responsive. The way that this was achieved was by first putting both parts of the program onto their own threads and then analysing what resources were being shared between the two threads. Anything that was being shared was then given a mutex to lock it so that all outputs of the code would still remain deterministic.

In theory this should have preserved functionality, and it did for the most part. However, bugs still worked their way in and due to the nature of multi-threaded applications, they were incredibly hard to track down and fix. But that's perhaps not the worst part is that performance didn't get better–it actually got worse. Herein lies the crux of the problem with multi-threading: it's absolutely a powerful tool, but not the right choice for every application. The reason for this is while adding a new thread to your application absolutely frees up other threads to work, by virtue of how you have to ensure thread safety you may actually end up with the same issue I have had. This is because when adding a mutex, there is an associated overhead with managing it. To add to this, if you're not careful about using it only as needed then you may block other threads enough to defeat the purpose of the multiple threads. As a final nail in the coffin, every thread has the problem of overhead. In order to create and manage a thread there is an amount of memory and CPU time required. If you cannot offset this overhead with the performance gained from moving functionality to a new thread, then you're always going to have worse performance.

So ultimately, the lesson to be learned here is that multi-threading is not a pancea. One must ensure they properly assess the benefits and drawbacks to multi-threading their application lest they end up with worse performance and more bugs than before.

2.2.4   Raycaster

Now for the moment (I hope) you've all been waiting for–raycasting! I need to get this out of the way before I say too much more: I *love* raycasters. I think they're some of the coolest ways we "hacked" graphics in the days of yore, and as such you'll probably find me gushing fairly frequently about just how cool I think they are. Sorry about that in advance.

So I suppose the next thing to explain is what raycasters are, and why they exist. In simple terms, raycasters are a way for us to cheat perspective to create 3D-like environments. To explain how it achieves this, we should explore how graphics works at a high level first.

In modern graphics APIs (such as OpenGL) we have what is called a graphics pipeline. This pipeline is made up of individual steps that we pass data through in order to process it into the pretty visuals we see on our screen. Each of these steps contains discrete programs which we call shaders, and of note to us is a specific shader called the fragment shader. Once all the processing on vertices, normals, etc is done by previous shaders we are now ready to render our fragments. For our purposes you can simply think of a fragment as a pixel

on your screen. Now imagine that for each fragment, we need to perform a calculation–on a 1920x1080 screen (a standard full HD display resolution) this means you have to make 2,073,600 calculations. To keep your computer feeling responsive you need to make sure that these calculations are done multiple times a second, which could bring up the calculations per second by a factor of 60. This is a lot of calculations! But our modern PCs with their very, very fancy GPUs are able to perform many of these calculations in parallel very, very quickly which allows our CPU to do other stuff while the PC still looks and feels responsive.

Old PCs weren't like this however. If they did have a GPU, it certainly wasn't able to perform this many discrete calculations at such a speed as to make true 3D graphics feasible. The raycaster gets around this with a very clever trick: only perform discrete calculations for every vertical scanline. This means that on our 1920x1080 screen, instead of you doing some 2 million calculations you are now only doing 1920–significantly less! So how does it achieve this? Well, buckle up because there's a fair bit of maths involved. Don't worry though! I'm really bad at maths so I'll be explaining things in simple terms as much as I can, since that's how I understand them (at a high level, anyway).

The first thing we start with is the top down 2D projection of the map that we wish to project with 3D perspective. This becomes our game world, and is actually what one might consider as the "real" world. This makes operating on entities in game incredibly easy since they're all entirely within 2D space. Next we place the player into the world and give them a viewing angle and cone. This cone will define where all raycasts are made. Now for the magic: we take the two extremes of the cone and through linear interpolation cast exactly as many rays as there are pixels horizontally in our render target. For each of these rays we calculate the distance between the player and the raycast hit. Depending on how far the hit was, we now draw a vertical bar on the render target at the appropriate vertical scanline–the closer the hit the larger the bar, and the further away the hit the smaller the bar. What we've now created is a really good, simple 3D projection that is incredibly performant and easy to work with.

Now if we want to add entities to the game, it's more or less the same idea except now we will be simply calculating the distance between each enemy and the player. The further away the entity is, the smaller the sprite drawn will end up being. And that's the bulk of the actual code! There are some more things to it such as depth testing (which will be talked about later) and lens correction which I won't really talk about here if only for the sake of brevity, but suffice to say that this is an incredibly clever and high performance way to create 3D graphics in games, during a time where such graphics were basically unheard of.

3   Challenges

This has been an absolutely enormous project, one that has spanned months at the time of writing this. That of course means that there have been many

different and interesting challenges that have arisen during development. I have not mentioned these challenges above for the sake of brevity, but for the interested reader I have detailed what they were, why they were problems, and how I did or didn't manage to resolve them.

3.1 Raycaster Depth Testing

Depth testing is a very, very large topic in graphics programming but is still generally simple enough. The idea is that when we draw something to the screen, we also want to write a value to the depth buffer. This value is simply how far whatever was drawn at that pixel is from the camera. What this allows us to do is perform a depth test on all subsequent draws to see if what we are trying to draw is in front of or behind something we just drew. This allows us to ensure that whenever we draw things to the screen, something that is behind another thing will not get drawn on top.

Since this is a 3D application, there is a simple depth buffer implemented. Recall that raycasters work entirely on vertical scanlines, which means that our depth buffer will also work on these vertical scanlines.The process of using this depth buffer goes as follows:

- Draw the walls in the scene, writing it's distance to the depth buffer

- Draw each entity in the scene, culling scanlines if a depth test shows it behind a wall

This is imperfect as a solution however, and that's because entities may be drawn on top of each other. You might assume that we can simply write to the depth buffer as we have already done for the walls, and this does actually work! But now we have a new kind of issue: transparency. Textures and sprites are not columns of a single colour, so when we draw them we need to sample them. Wall textures are fairly trivial in this regard since we know that they're always going to be solid with no transparency, so we always draw every pixel of the texture. Sprites for the entities are a little bit more complicated however, and this is because they can and do include transparency. Now, the way transparency is handled in this very simple (in graphics APIs like OpenGL it can get far more complicated with blend modes and render queues) since all we do is check if something is transparent at all or not, no translucency is handled. However, this transparency testing is only done during the sampling of the sprite, there is no information about it written into the depth buffer.

We could attempt to ensure that all sprites are written to the depth buffer, and that any transparency is ignored and not written to the depth buffer. This only works to a point however, because if you recall the depth buffer is only tracking values for every scanline and there are absolutely cases where a scanline can be partially transparent. There are many methods to achieve this kind of depth testing to ensure that sprites are rendered in the correct order, but I have opted for a brute force method: rendering the entities in ascending order of distance. This means that the entities which are furthest away from the

player will get rendered first and the closest will be rendered last. This works simply because anything in front of other entities will just render over it. It's of course not ideal–we're running a sorting algorithm every frame and we're wasting resources on rendering more than we can see. A lot of CPU time gets wasted this way, but it's easy to implement and absolutely works brilliantly.

3.2 Mouse Looking

With the way that the raycaster had been written, changing the player's view angle would be trivial–and indeed it is. There is a simple method on the Engine class that allows you to adjust the current player's view angle, which is stored in radians, by a provided delta. At first this was achieved with a simple button press which was easy to create: while the key is held down, increment the view angle by a set delta. This worked excellently, but it would be *way* cooler if you could control it with your mouse instead.

This is where the complications begin, but not because mouse movements are hard to track–quite the contrary actually! Really, all that needed to be done was store the previous position of the mouse, store the next position of the mouse, and then calculate the magnitude of the two vectors. In fact, this is made even easier since the only axis we care about is the x axis (raycasters by design only allow for looking left and right) so we need only store the previous and current x axis and calculate the difference between the two to get our delta.

In SFML we have, much like SDL does, an event loop. These are pretty simple, allowing us to listen for all window events. It works incredibly well, and indeed is perfect for most of what is needed in this project. So when handling a mouse event, it would seem most reasonable to also add it to this event loop. In fact, checking the SFML documentation we can see that there is actually an event type called MouseMoveEvent which is perfect for us. So in theory, we should just be able to add this event to the loop and calculate our deltas like this:

```
sf::Event event;
int last_x;
while ( window->pollEvent( event ) ) {
    // handle other event types here

    switch ( event.type ) {
        case sf::Event::MouseMoveEvent:
            const int delta = event.mouseMove.x - last_x;
            last_x = event.mouseMove.x;
            move_view( delta );
            break;
    }
}
```

Here is where we encounter our first problem, since if we just calculate the deltas between the last position and the current position then we may eventually

(and probably will) run out of screen space to travel. To fix this is actually fairly simple, since we need only do the following:

```
// on program init
window->setMouseCursorGrabbed( true );

// during event loop
sf::Mouse::setPosition( window_center, *window );
```

The general idea is that we are ensuring that the cursor is locked to the window so we can always capture events, and we're resetting the position to the center of the window every trip through the loop. This works well, but introduces yet another issue. You see, when we set the mouse position it actually adds another MouseMoveEvent to the event stack. This is a problem, because we're in the middle of polling that very same stack and if we add another event to that stack while we're in the event loop... Well you're just going to spin endlessly in a loop–which is exactly what ended up happening.

This problem was the cause for much pain over the course of a few days as I couldn't figure out how to stop the event stack from registering a new event when I move the mouse cursor through code. Eventually I happened across a post on the SFML forums (which unfortunately, I have since lost) which keyed me in to the solution which I have ultimately landed on. That solution is a simple one: don't use the event loop. It might be a hack, it might actually be intended, really I'm not actually sure. What I do know, however, is that calculating the delta myself outside of the event loop ended up being the solution I needed to side step the issues with endless looping. However, I do fully admit that there's a potential issue with performance from this approach (perhaps a small one, but and issue all the same) since we're still adding these events to the stack and we're now doing delta calculations on every single frame. These are sacrifices I'm willing to make however since early optimisation is the enemy of progress, and there is only so much time to complete this project. Since it didn't pose any immediate problem with the performance of the project, I've left it like this in the final product.

4   Conclusion